

Ethics of AI

CASE BRIEF

X's Grok included an image tool capable of producing manipulated, sexualized images of real people, including minors, without their consent. In January 2026 the European Commission opened a formal Digital Services Act (DSA) investigation into X, examining whether the company properly assessed and mitigated the risks of Grok before deployment, and whether X's recommender systems then amplified the resulting content. Commission President von der Leyen said Europe would not tolerate the "digital undressing" of women and children. Separately, at least eight countries pursued their own regulatory actions against X or xAI over the same underlying problem, but through different legal tools — data protection law, criminal law, child-safety law, platform regulation — rather than one coordinated instrument.

INDIVIDUAL SORTING

Sort each item below into: (a) non-maleficence norm, (b) beneficence norm, (c) unclear/both. Be ready to justify.

1. A legal requirement that platforms conduct a risk assessment before deploying a high-risk AI feature.
2. A legal requirement to mitigate identified risks after they're found (e.g., restrict or filter the tool's outputs).
3. A country's criminal law used to prosecute the distribution of abusive imagery generated through the tool.
4. A design decision to detect and block an attempt to generate a sexualized image of an identifiable real person before the image is ever generated.
5. A public company statement promising "we care about user safety," with no enforceable mechanism behind it.
6. A requirement that platforms design tools to actively protect the dignity of vulnerable groups, not merely avoid the worst-case outcome.
7. Age-verification requirements before granting access to a generative image tool.
8. A public awareness campaign teaching users how to report abusive deepfakes.

DISCUSSION

The DSA's risk-assessment and mitigation obligations apply only after a company has built and shipped the feature — enforcement occurs after the risk is identified. Is a norm that only acts after the fact still doing the work of beneficence, or is it just a stricter version of non-maleficence? Where would a norm have to sit to actually count as an instance of beneficence?